AMOGH JAGADISH TAMBAD

(480) 876-5096 | San Francisco, CA | tambadamogh@gmail.com | linkedin.com/in/ajtambad | github.com/Ajtambad

EDUCATION

Master of Science, Computer Science May 2025
Arizona State University, Tempe, AZ 4.00 GPA
Coursework: Cloud Computing, Data Processing at Scale, Data Mining, Software Security

Bachelor of Technology, Computer Science May 2021
REVA University, Bangalore, India 3.77 GPA
Coursework: Data Structure and Algorithms, Operating Systems

SKILLS

Programming Languages: Python, C++ Go, Java, JavaScript, Typescript, Scala, HTML, XML
Technologies: AWS (EC2, Lambda, S3, ECS, CloudFormation, EKS, ECR, CloudWatch), GCP, TensorFlow, Docker, Kubernetes, Terraform, Knative, Jenkins, Github Actions, React, Angular, Flask, Node.js, GCP
Tools: Linux (RHEL, Ubuntu), Unix, Git (Version Control), Bash, Github, Ansible, Chef, Gitlab, Nginx, Kafka, Redis, Prometheus, Grafana, Postman, Cursor, CoPilot, Claude Code
Database Systems: SQL(MySQL, PostgreSQL), NoSQL (MongoDB)
LLM/GenAI Frameworks: MCP, LangChain, LangGraph, RAG, VectorDBs (FAISS, ChromaDB)
Miscellaneous: Distributed Systems, RESTful APIs, Microservices, Object Oriented Programming, Agile, Test Driven Development (TDD), Infrastructure as code (IaC), DevOps, Cloud Infrastructure, Communication Skills, Collaboration, Writing

PROFESSIONAL EXPERIENCE

Research Assistant, VISA Lab Jun 2025 - Present
Arizona State University, Remote

- Developed **FlowBench**, a workflow-based distributed benchmark by leveraging **Python**, **Docker**, **Kubernetes**, and **FaaS** principles to evaluate custom edge computing applications, providing a comprehensive report on 6+ metrics.
- Built and tested a video analytics workflow via **OpenCV** on a containerized microservices architecture with Kubernetes, processing 10,000+ frames per minute.

Site Reliability Engineer (SRE) Intern Jun 2024 - Aug 2024
Arch Mortgage Insurance, Greensboro, NC

- Filtered and routed logs from OpenShift Kubernetes clusters to Splunk using Cribl Stream pipelines, reducing daily Splunk data ingestion by **40–50 GB** and improving **log search performance by 20%**
- Designed an automated pipeline using Ansible and Red Hat registry APIs to sync updated catalog images to Nexus Repository, reducing manual update time by **90%**

System Engineer - 1 May 2021 - Jul 2023
Cerner Healthcare (Oracle Health), Bangalore, India

- Migrated 80% of data from on-prem to **AWS**, enhancing data access flexibility, security, and cost-efficiency
- Participated in regular on-call rotations with cross-functional teams, leveraging **Zabbix** and **Splunk** for system health monitoring, quality assurance, and resolving production alerts, maintaining **99.99%** service reliability
- Troubleshot and performed root cause analysis **Jenkins** pipeline issues, minimizing support ticket resolution time by **40%** and ensuring **99.9%** uptime for **CI/CD** workflows, leading to uninterrupted deployment pipelines
- Managed **300+** bi-weekly microservice deployments, including Splunk and non-Splunk-based services, using **Chef**, while also tracking them with **JIRA**, accelerating delivery of new UI and backend features in a fast-paced production environment

ACADEMIC PROJECTS

End-to-End Deployment Automation Apr 2025

- Automated end-to-end AWS EC2 provisioning and cloud infrastructure management using **Terraform**, **Ansible**, **Jenkins**, and **GitHub Actions**, enabling infrastructure setup and hands-free web service deployments based on version control changes.
- Built and optimized **CI/CD** pipelines to dynamically retrieve instance IPs, configure secure SSH access, and deploy services, eliminating manual intervention and resolving **IAM** and resource issues in production-like environments

AWS-Based Face Recognition App May 2024

- Developed and deployed a **Flask**-based image recognition app using **Gunicorn** on AWS EC2, enabling HTTP REST API based uploads and forwarding images to **S3** via **SQS** for asynchronous processing using AWS service API.
- Designed an auto-scaling app tier that scaled up to 20 EC2 instances based on **SQS** queue depth, ensuring efficient, real-time image processing under varying workloads